

# CS 214 Homework 4

Fall 2021

## Introduction

In this assignment, you'll implement a custom version of malloc, free, and realloc. You should implement these in `mymalloc.c` and `mymalloc.h`. These files should not contain a `main` function, so you may want a separate `.c` file for testing. You'll also write a performance testing program in `memperf.c`.

## Library functionality

You should implement these functions, described below:

```
void    myinit(int allocAlg);
void*   mymalloc(size_t size);
void    myfree(void* ptr);
void*   myrealloc(void* ptr, size_t size);
void    mycleanup();
double  utilization();
```

- `myinit(allocAlg)`

Create a 1 MB “heap” and perform any other initializations your code needs. You can assume any application using your library will call this first.

The `allocAlg` argument describes what algorithm to use to find a free block (see Lec. 17, p. 19):

- 0: first fit
- 1: next fit
- 2: best fit

- `mymalloc(size)`

From the “heap”, allocate a region of at least the requested `size` and return a pointer to the beginning of the region. If it cannot be allocated, return `NULL`.

All returned addresses must be 8-byte aligned. That is, the region you allocate should start at an address that's divisible by 8.

If `size` is 0, `mymalloc` does nothing and returns `NULL`.

- `myfree(ptr)`

Mark the given region as free and available to be allocated for future requests. It should be coalesced with adjacent free regions.

You should maintain an explicit free list.

If `ptr` is `NULL`, `myfree` does nothing.

You should detect a number of error conditions:

- trying to free an invalid address

If the address given isn't in your "heap", print "error: not a heap pointer". (For example, if your "heap" ran from 0x4000 to 0x6000, but the `ptr` given to `myfree` was 0x10.)

If the address given is in your "heap" but wasn't returned by `mymalloc`, print "error: not a malloced address". (For example, if `mymalloc` returns 0x4040, but the `ptr` given to `myfree` was 0x4041.)

- double free

If the address given was returned by `mymalloc` but has already been freed, print "error: double free".

- `myrealloc(ptr, size)`

Reallocate the region pointed to by `ptr` to be at least the new given `size`. If this cannot be done in-place, a new region should be allocated, the data from the original region should be copied over, and the old region should be freed.

If the reallocation can't be done, return `NULL`.

If `ptr` is `NULL`, this is equivalent to `mymalloc(size)`.

If `size` is 0, this is equivalent to `myfree(ptr)` and `myrealloc` returns `NULL`.

If both `ptr` is `NULL` and `size` is 0, `myrealloc` does nothing and returns `NULL`.

- `mycleanup()`

Free the 1 MB "heap" and perform any other cleanup your code needs. You can assume any application using your library will call this last.

Your library should support "resetting" everything by calling `mycleanup` followed by `myinit`, since the performance testing program will need to do this.

- `utilization()`

Calculate space utilization as a ratio of memory used vs. space used in the "heap". For example, if you called `mymalloc(50)` followed by `mymalloc(64)` and determined that you needed 128 bytes on the "heap" to satisfy these requests, the ratio would be

$$(50 + 64)/128 \approx 0.89.$$

Blocks that have been freed should not count towards memory used. The space used in the "heap" is the distance from the beginning of the "heap" to the end of the last allocated block.

## C standard library usage

Outside of `myinit` and `mycleanup`, you should not call the standard `malloc`, `calloc`, `free`, or `realloc` functions.

## Performance

You should write another program, `memperf.c`, that uses your library and attempts to measure its performance. It should randomly allocate, reallocate, and free blocks, and then analyze throughput and utilization for each allocation algorithm (call `mycleanup` to finish testing one algorithm and `myinit` again to start another). Blocks should be randomly sized, say from 1 byte to 256 bytes.

Utilization is defined above in the description of the `utilization()` function. Throughput should be measured as number of `mymalloc` / `myrealloc` / `myfree` operations per second. You may need to tweak the number of operations and block size to get reliable results without `memperf` taking too long to run or exhausting the available memory in your “heap”. About 1,000,000 operations may be a good starting point.

Some possibly useful functions: `rand`, `srand`, `time`, `gettimeofday`.

The output format should be as follows:

```
$ ./memperf
First fit throughput: 7812.5 ops/sec
First fit utilization: 0.85
Next fit throughput: 15625 ops/sec
Next fit utilization: 0.9
Best fit throughput: 3906.25 ops/sec
Best fit utilization: 0.92
```

## Compiling and testing

You should create a Makefile so that running `make` or `make all` builds your program (i.e., `memperf`). You should use similar `CFLAGS` to `gcc` as in previous homeworks, e.g.:

```
-g -Wall -Wvla -fsanitize=address
```

## Submission

If you develop on your local machine, please be sure to test your code on ilab before submitting.

Please submit the assignment on Canvas as a tar file `hw4.tar` that, when expanded, produces a `hw4` directory (possibly with additional `.c/.h` files):

```
hw4
├── Makefile
├── memperf.c
├── mymalloc.c
└── mymalloc.h
```

If you work in a group, please include a `README.txt` file with the names and netIDs of everyone in your group. Only one person needs to submit on Canvas.